

PostgreSQL Indexes

Overview

Boost read performance:

Efficient data structures that speed up data retrieval.

Reduce disk I/O:

By scanning less data, queries become faster.

Trade-off: Slower writes

Indexes need to be updated on INSERT, UPDATE, DELETE operations, which can impact write performance.

Index Types

B-TREE
(default)

GIN

HASH

GiST

Data Model

transactions	
🔑	id integer
🔑	account_id integer
⌚	transaction_type character varying(20)
⌚	amount numeric(15,2)
⌚	transaction_date timestamp without time zone
🔑	target_account_id integer

accounts	
🔑	id integer
🔑	customer_id integer
⌚	account_type character varying(20)
⌚	balance numeric(15,2)
⌚	created_at timestamp without time zone

customers	
🔑	id integer
⌚	first_name character varying(50)
⌚	last_name character varying(50)
❶	email character varying(100)
⌚	phone character varying(20)
⌚	address text
⌚	created_at timestamp without time zone

Query without index

```
EXPLAIN ANALYZE SELECT * FROM transactions WHERE amount > 500;
```

QUERY PLAN

text

```
Gather (cost=1000.00..13785.33 rows=100 width=34) (actual time=62.354..66.524 rows=0 loops=1)
```

Workers Planned: 2

Workers Launched: 2

```
-> Parallel Seq Scan on transactions (cost=0.00..12775.33 rows=42 width=34) (actual time=58.934..58.935 rows=0 loop.
```

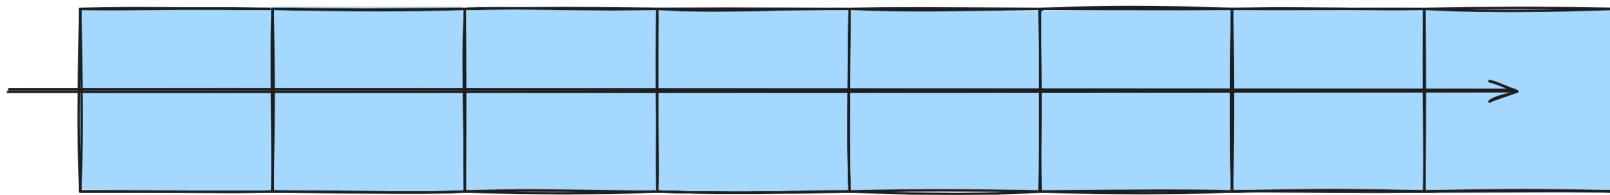
Filter: (amount > '500'::numeric)

Rows Removed by Filter: 333333

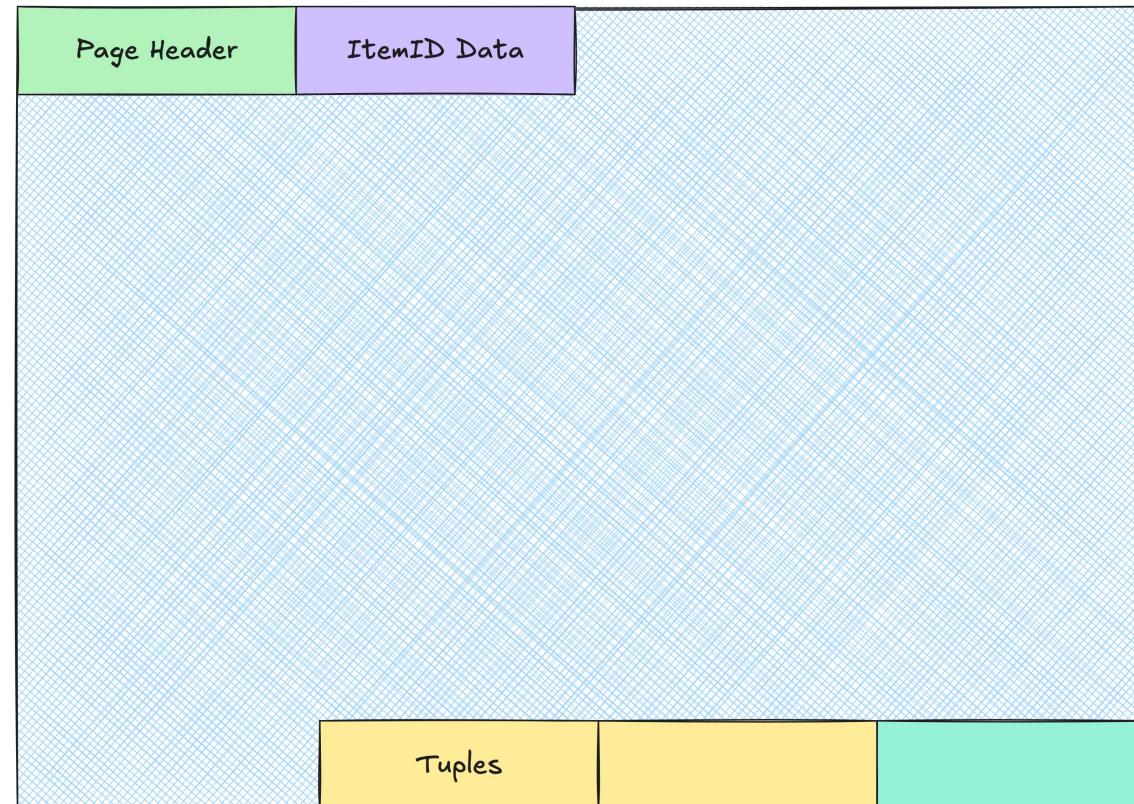
Planning Time: 0.074 ms

Execution Time: 66.645 ms

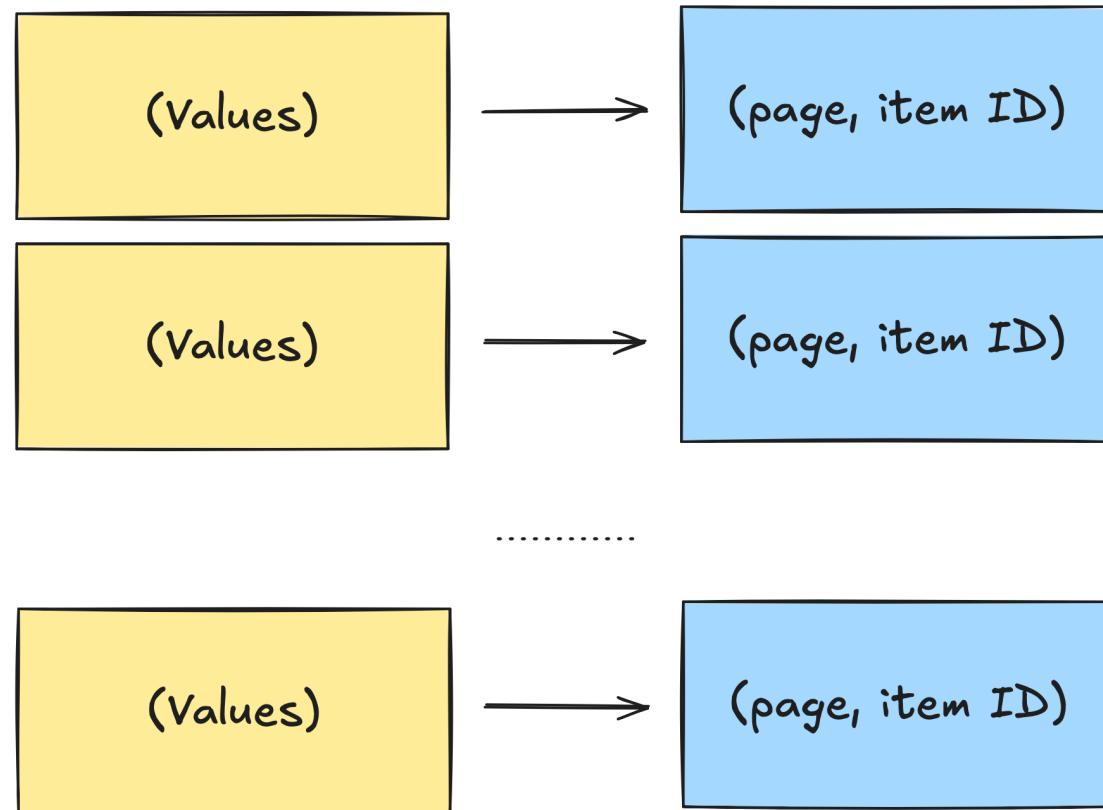
Full Table Scan



How Postgres stores data

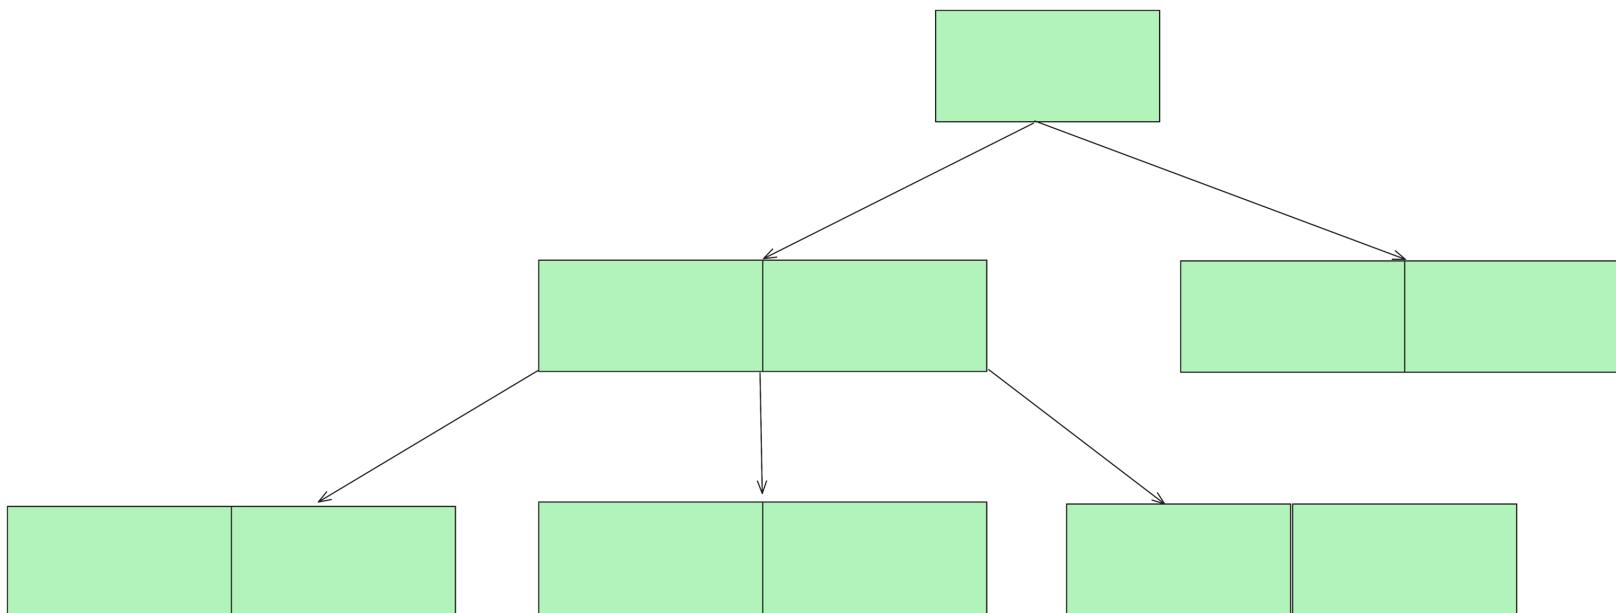


What is an index

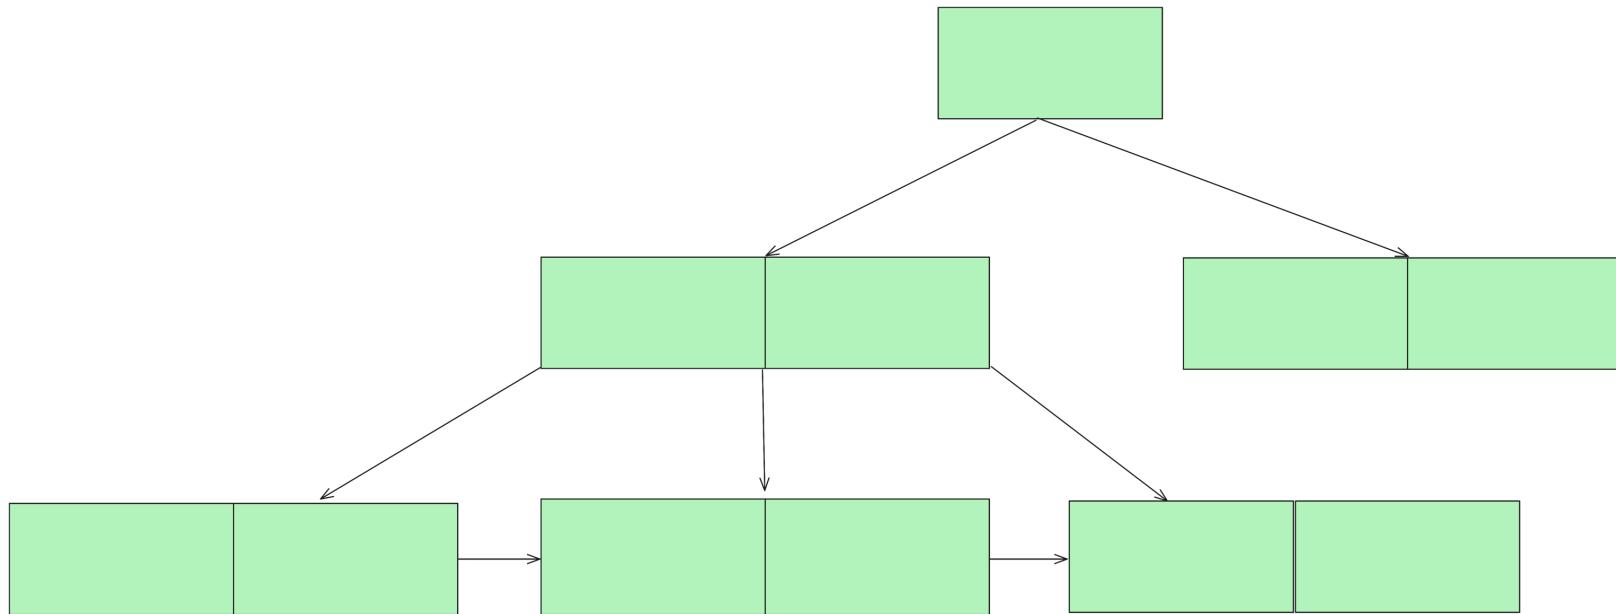


B-TREE

Self-balancing search tree



B+TREE



Query with index

```
SELECT * FROM transactions WHERE amount > 500;
```

```
CREATE INDEX transaction_amount_idx ON transactions(amount);
```

QUERY PLAN

text

Index Scan using transaction_amount_idx on transactions (cost=0.42..8.44 rows=1 width=34) (actual time=0.005..0.005 rows=0 loop..)

Index Cond: (amount > '500'::numeric)

Planning Time: 0.320 ms

Execution Time: 0.017 ms was 66.645 ms

Composite index

```
EXPLAIN ANALYZE SELECT amount, transaction_date
  FROM transactions
 WHERE amount > 500 AND transaction_date > NOW() - INTERVAL '10 days';
```

QUERY PLAN

text

Gather (cost=1000.00..16903.23 rows=29 width=14) (actual time=80.116..85.207 rows=0 loops=1)

Workers Planned: 2

Workers Launched: 2

-> Parallel Seq Scan on transactions (cost=0.00..15900.33 rows=12 width=14) (actual time=76.336..76.336 rows=0)

Filter: ((amount > '500'::numeric) AND (transaction_date > (now() - '10 days'::interval)))

Rows Removed by Filter: 333333

Planning Time: 0.166 ms

Execution Time: 85.225 ms

Composite index

```
EXPLAIN ANALYZE SELECT amount, transaction_date  
  FROM transactions  
 WHERE amount > 500 AND transaction_date > NOW() - INTERVAL '10 days';
```

```
CREATE INDEX transaction_amount_and_date_idx ON transactions(amount, transaction_date);
```

QUERY PLAN

text

Index Only Scan using transaction_amount_and_date_idx on transactions (cost=0.43..4.45 rows=1 width=14) (actual time=0.007..0.008 rows=0)

Index Cond: ((amount > '500'::numeric) AND (transaction_date > (now() - '10 days'::interval)))

Heap Fetches: 0

Planning Time: 0.486 ms

Execution Time: 0.021 ms was 85.225 ms

Composite index

```
EXPLAIN ANALYZE SELECT amount, transaction_date  
  FROM transactions  
 WHERE transaction_date > NOW() - INTERVAL '10 days';
```

```
CREATE INDEX transaction_amount_and_date_idx ON transactions(amount, transaction_date);
```

QUERY PLAN

text

Seq Scan on transactions (cost=0.00..25067.00 rows=290884 width=14) (actual time=0.013..381.238 rows=290714)

Filter: (transaction_date > (now() - '10 days'::interval))

Rows Removed by Filter: 709286

Planning Time: 0.062 ms

Execution Time: 396.937 ms

Partial Index

```
CREATE INDEX transaction_status_idx ON transactions(amount)
WHERE transaction_status = 'in-progress';
```

"in-progress" – 10%

Reduce the size of the index

Covering Index

```
SELECT account_id, transaction_type  
FROM transactions WHERE amount > 150;
```

```
CREATE INDEX transaction_amount_idx ON transactions(amount);
```

Bitmap Heap Scan on transactions (cost=2447.47..11645.18 rows=130457 width=12) (actual time=29.154..74.158 rows=134064 loops=1)
Recheck Cond: (amount > '150'::numeric)
Heap Blocks: exact=7567
-> Bitmap Index Scan on transaction_amount_idx (cost=0.00..2414.85 rows=130457 width=0) (actual time=27.532..27.533 rows=134064 loop...
Index Cond: (amount > '150'::numeric)
Planning Time: 0.219 ms
Execution Time: 81.909 ms

Covering Index

```
SELECT account_id, transaction_type  
FROM transactions WHERE amount > 150;
```

```
CREATE INDEX transaction_amount_covering_idx ON transactions(amount)  
INCLUDE (account_id, transaction_type);
```

Index Only Scan using transaction_amount_covering_idx on transactions (cost=0.42..4871.42 rows=130457 width=12)

Index Cond: (amount > '150'::numeric)

Heap Fetches: 0

Planning Time: 0.085 ms

Execution Time: 38.894 ms

Index on Join Column

```
SELECT c.first_name, c.last_name, a.created_at FROM customers c
JOIN accounts a ON a.customer_id = c.id
WHERE c.id = 42;
```

QUERY PLAN

text

Nested Loop (cost=0.29..670.98 rows=1 width=21) (actual time=0.025..2.434 rows=1 loops=1)

-> Index Scan using customers_pkey on customers c (cost=0.29..8.31 rows=1 width=17) (actual time=0.012..0.014 rows=1)

Index Cond: (id = 42)

-> Seq Scan on accounts a (cost=0.00..662.66 rows=1 width=12) (actual time=0.011..2.417 rows=1 loops=1)

Filter: (customer_id = 42)

Rows Removed by Filter: 33332

Planning Time: 0.220 ms

Execution Time: 2.456 ms

Index on Join Column

```
SELECT c.first_name, c.last_name, a.created_at FROM customers c
JOIN accounts a ON a.customer_id = c.id
WHERE c.id = 42;
```

```
CREATE INDEX accounts_customer_id_idx ON accounts(customer_id);
```

QUERY PLAN

text

Nested Loop (cost=0.58..16.62 rows=1 width=21) (actual time=0.034..0.035 rows=1 loops=1)

-> Index Scan using customers_pkey on customers c (cost=0.29..8.31 rows=1 width=17) (actual time=0.009..0.009 rows=1 loops=1)

 Index Cond: (id = 42)

-> Index Scan using accounts_customer_id_idx on accounts a (cost=0.29..8.31 rows=1 width=12) (actual time=0.022..0.022 rows=1)

 Index Cond: (customer_id = 42)

Planning Time: 1.463 ms

Execution Time: 0.054 ms

Function Based Index

```
CREATE INDEX idx_customers_lower_first_name ON customers(email);
```

```
SELECT first_name FROM customers WHERE lower(email) = 'test@example.com';
```

QUERY PLAN

text

```
Seq Scan on customers (cost=0.00..1052.99 rows=167 width=6) (actual time=38.595..38.596 rows=0 |
```

```
    Filter: (lower((email)::text) = 'test@example.com'::text)
```

```
    Rows Removed by Filter: 33333
```

```
Planning Time: 0.338 ms
```

```
Execution Time: 38.625 ms
```

```
CREATE INDEX idx_customers_lower_first_name ON customers(lower(email));
```

Query With Casting

```
SELECT * FROM customers WHERE CAST(created_at AS DATE) = '2025-12-13';
```

```
CREATE INDEX idx_customers_date ON customers(created_at);
```

Seq Scan on customers (cost=0.00..1052.99 rows=167 width=101) (actual time=4.758..4.758 rows=0 loops...)

Filter: ((created_at)::date = '2025-12-13'::date)

Rows Removed by Filter: 33333

Planning Time: 0.051 ms

Execution Time: 4.772 ms

```
CREATE INDEX idx_customers_date ON customers((CAST(created_at AS DATE)));
```

Query With Like

```
SELECT * FROM customers WHERE first_name LIKE 'sa%';
```

```
CREATE INDEX idx_customers_first_name ON customers(first_name);
```

Seq Scan on customers (cost=0.00..969.66 rows=1 width=101) (actual time=2.984..2.984)

Filter: ((first_name)::text ~~ 'sa%'::text)

Rows Removed by Filter: 33333

Planning Time: 1.319 ms

Execution Time: 2.996 ms

```
CREATE INDEX idx_customers_first_name ON customers(first_name text_pattern_ops);
```

Low Selectivity

```
SELECT * FROM transactions  
WHERE transaction_type = 'deposit';
```

The transaction type 'deposit' has a frequency of 80% among all transactions.

```
CREATE INDEX idx_transactions_type ON transactions(transaction_type);
```

Seq Scan on transactions (cost=0.00..20067.00 rows=798667 width=34) (actual time=0.007..118.292 rows=799830)

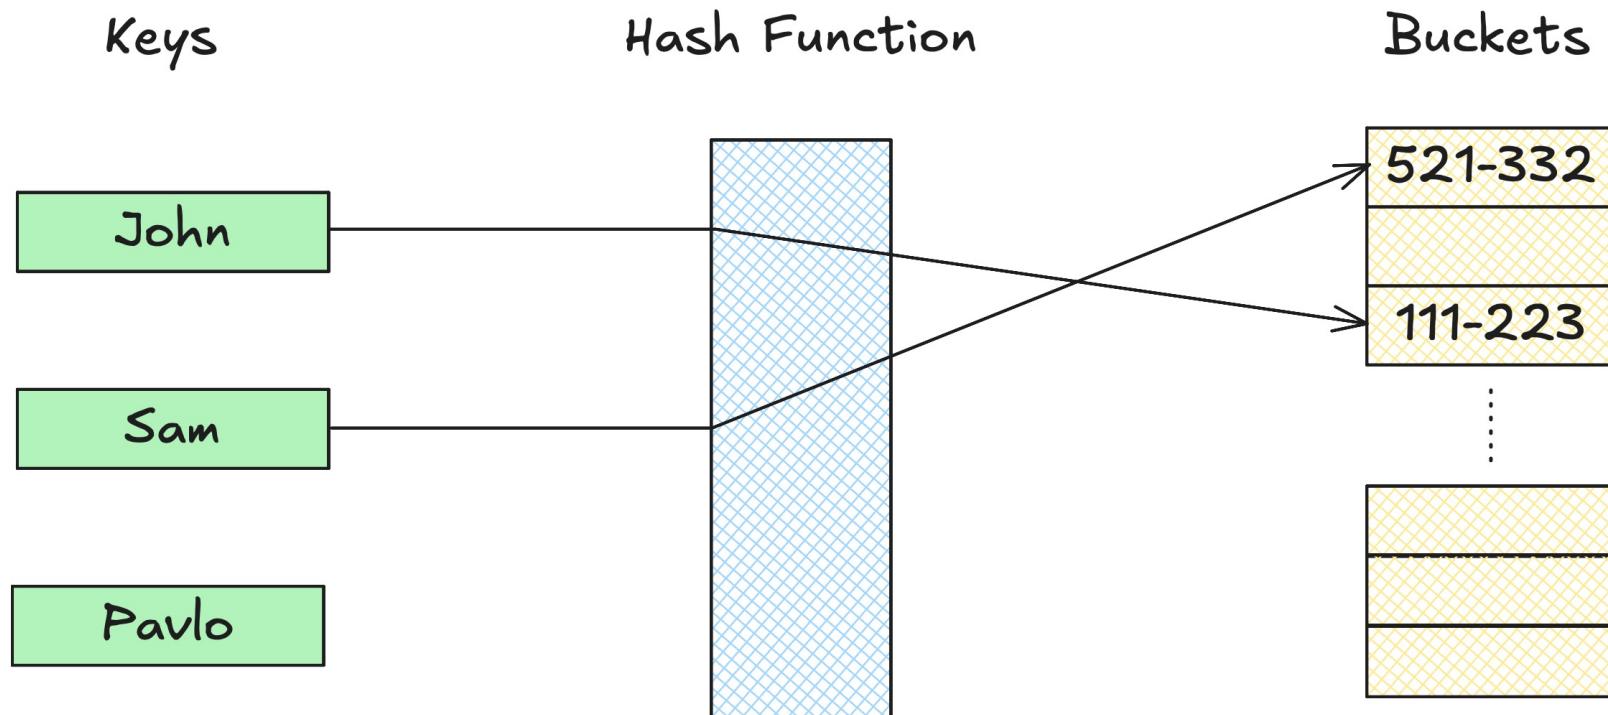
Filter: ((transaction_type)::text = 'deposit'::text)

Rows Removed by Filter: 200170

Planning Time: 0.227 ms

Execution Time: 148.834 ms

Hash Index



Hash Index

```
SELECT first_name FROM customers WHERE first_name = 'David';
```

```
CREATE INDEX idx_customers_first_name_hash ON customers USING hash (first_name);
```

Index Scan on idx_customers_first_name_hash

IMPORTANT: Hash index works **only** with the = operator.

Thank you

- Author: Dmytro Halchenko
- My LinkedIn: <https://www.linkedin.com/in/dmytro-halchenko-0307751b6/>
- Date: March 2025
- [Join Codeus community in Discord](#)
- [Join Codeus community in LinkedIn](#)